

2027 UF Guide to AI in Science

Nicholas A. Coles Abhay Alaukik Majed A. Arishi
Shanting Chen Annabel V. Dang Jungjae Im Nabiha Iqbal
Minyi Jiang Rui Jin Zhongchi Li Sooyeon Park
Mariana G. Rebello Jihyun Soh Paola J. Sullivan
Lyra G. Toohey Sofia C. Villanueva

2026-09-08

Table of contents

Ideation and Project Management	6
Ideation and Project Management: Using AI	7
Areas of consensus — and dissensus	8
Ideation and Project Management: Critiquing AI	10
Areas of consensus — and dissensus	10
 Literature Review and Synthesis	 13
Literature Review and Synthesis: Using AI	14
Literature Review and Synthesis: Critiquing AI	15
 Data Collection and Generation	 16
Data Collection and Generation: Using AI	17
Data Collection and Generation: Critiquing AI	18
 Analysis and Modeling	 19
Analysis and Modeling: Using AI	20
Analysis and Modeling: Critiquing AI	21
 Writing and Communication	 22
Writing and Communication: Using AI	23
Writing and Communication: Critiquing AI	24

Ethics and Outreach	25
Ethics and Outreach: Using AI	26
Ethics and Outreach: Critiquing AI	27
Personal AI Philosophy Statements	28
References	29

This was written by humans.

But we provide links to helpful AI-generated content along the way.

AI is radically redefining the nature of scientific inquiry — for better and for worse.

In Fall 2026, faculty, graduate students, and postdocs across multiple departments at the University of Florida came together to understand this urgent development. In a seminar titled *Collective [Artificial] Intelligence in Science*, we documented and discussed how AI is changing each component of our research workflow. This includes project management, idea development, data analysis, writing, and dissemination. To push our understanding, we alternated between the perspectives of an AI zealot and an AI critic — identifying key resources, consensus statements, and areas of disagreement along the way. The result is the **2027 UF Guide to AI in Science**.

The seminar began with an introduction to big-team science — its rise¹, competitive advantages², and relevance to the current moment. For instance, the paper describing OpenAI’s GPT-1 had 4 authors³; the paper describing GPT-3 had 31 authors⁴; and the paper describing GPT-5 had 483 authors⁵ (Figure 1). Similarly, a prominent AI-led cybersecurity attack on the online community, Hugging Face, involved coordinated action among not one but *thousands* of AI agents⁶.




Figure 1: Credited co-authors on flagship AI model reports from OpenAI, Google, and Meta, 2018–2025. One square is one credited person; each block’s area is its author list.

After reviewing these developments, the course itself became a large-scale collaboration, dedicated to understanding how AI is changing the nature of science. Each week followed a consistent structure (Figure 2). First, we would select a component of the research pipeline to examine (e.g., data analysis). Second, we would spend one week compiling resources for understanding the emerging role of AI in that part of the research pipeline. Third, we would engage in peer-to-peer teaching. Last, we would use a live online voting platform to solicit, debate, and vote upon key resources and takeaways for the week^{7,8}. Through these exercises, seminar participants received hands on experience with large-scale collaboration^{9,10}. Along the way, not only developed and debated their *personal* philosophy statements – but also the general direction of AI at UF. Their results culminate into this **2027 UF Guide to AI in Science**: a summary of our top resources, takeaways, and disagreements.




Figure 2: Structure of the UF Fall 2026 seminar, titled Collective [Artificial] Intelligence in Science.

Ideation and Project Management


Ideation and Project Management: Using AI

This was written by humans.

But we also like this [AI-generated podcast](#) (20:36) of the entire 2-week unit on Ideation and Project Management.

Many researchers' first contact with AI tools is in the realm of ideation or project management (Figure 1). For example, an in-class demo focused on an AI-assisted workflow for managing research trainee projects. More specifically, one seminar participant has their trainees submit weekly updates to a folder shared with the LLM, Claude (Figure 2, Panel A). This allows them to later prompt the model with questions about past conversations, challenges, decisions, and solutions (Figure 2, Panel B). Specifically, the seminar participant demonstrated how Claude can be used to efficiently take stock of a trainee's progress, have a conversation with the LLM about potential future directions, and generate a simulated figure to send back to the trainee to convey next steps (Figure 2, Panel C).

Top discussion points



Top resources




Figure 1: Overview of Week 2 of the Fall 2026 Collective [Artificial] Intelligence Seminar. Left: top discussion points, along with estimated level of agreement among seminar participants. Right: three most popular resources identified by seminar participants.

The resources compiled by seminar participants illustrate that AI is a tool that is increasingly being used to expand ideation and project management capabilities (Figure 1). For example, as early as 2025, researchers documented that LLMs were capable of generating research ideas that were judged more novel than those of human experts¹¹. These capabilities can be further improved by integrating causal discovery algorithms. More specifically, LLMs can efficiently extract which causal relationships have been reported in past research. When mapped across studies, causal discovery algorithms can be used to identify gaps in causal network maps — i.e., potential future research directions¹². Such results highlight the increasing sophistication with which researchers are deploying AI to organize information, extract ideas, and even develop theory¹³.

Areas of consensus — and dissensus

Disagreement is common when researchers collaborate^{7,8} — and our seminar discussion was marked by both consensus and dissensus.

In terms of consensus, there was broad agreement that LLMs can assist with ideation, project management, and the organization of information. There was also broad agreement that this is changing the nature of competition in science — placing more emphasis on who can better leverage these tools for ideation and execution.

Yet, there was clear disagreement about how far things should go. Seminar participants strongly disagreed about whether AI is reliably superior at ideation, whether the institution of science should shift to more AI-led work, and whether humanity should strive for fully autonomous AI scientists. These disagreements echo throughout the scientific landscape, with recent back-to-back *Nature* pieces on both the capabilities¹⁴ and limitations¹⁵ of science with fewer humans and more AI. These limitations became the focus of the following week of the Fall 2026 Collective [Artificial] Intelligence Seminar (see Section).

[illegible]

```
Claude Code v2.1.245
Opus 5 · Claude Max
~/Documents/GitHub/LabManual/checkin/Cindy

> /model
  | Set model to Opus 5 (1M context) and saved as your default for new sessions

> run through cindy's notes and give me a sense for how she is identifying "DEI scholars"

• I'll look through Cindy's check-in notes.

  Listed 1 directory, ran 5 shell commands (ctrl+o to expand)

• Read all 19 of Cindy's check-ins (04/01-08/18). Here's how "DEI scholar" actually gets built, and where it wobbles.

  It's a two-stage construct: paper label → author score → cutoff
```

Texas

Line graph: % outmigrant (2018-2025). BF values: $BF_{1,0} = 4.2 \times 10^3$, $BF_{1,2} = 8.1 \times 10^3$, $BF_{1,3} = 2.6 \times 10^3$.

Florida

Line graph: % outmigrant (2018-2025). BF values: $BF_{1,0} = 1.8 \times 10^3$, $BF_{1,2} = 3.4 \times 10^3$, $BF_{1,3} = 170$.

Tennessee

Line graph: % outmigrant (2018-2025). BF values: $BF_{1,0} = 29$, $BF_{1,2} = 4.1$, $BF_{1,3} = 1.2$.

United States

Line graph: % outmigrant (2018-2025). BF values: $BF_{1,0} = 140$, $BF_{1,2} = 0.3$, $BF_{1,3} = 0.42$.

Legend: Texas (orange), Florida (green), Tennessee (pink), New York (dark blue), California (yellow), United States (dark blue), Canada (teal), United Kingdom (maroon).

9

Ideation and Project Management: Critiquing AI

This was written by humans.

But we also like this [AI-generated podcast](#) (20:36) of the entire 2-week unit on Ideation and Project Management.

In February 2026, an engineer running an online educational program used Claude Code to help launch a new side project. Instead of launching something new, they ultimately deleted everything they had already built: servers, networks, and years' worth of datasets. The failure point, the engineer would later suggest, was over-automation¹⁶. The engineer never checked whether backups could be restored; they never pushed back on questionable recommendations from the coding agent; and they allowed it to run irreversible commands completely unsupervised. Issues like these are bound to become more common as scientists increasingly rely on AI — the focus of Week 3 of the Fall 2026 Collective [Artificial] Intelligence Seminar (Figure 1).


This week's most popular resources focused on the pitfalls of using AI for ideation and project management (Figure 1). For example, one study used electroencephalography to measure brain connectivity patterns while writing alone, with a search engine, or with an LLM. Results suggested that LLM users exhibited the lowest levels of distributed brain activity — an index of lower cognitive engagement¹⁷. Similarly, when researchers rely more on AI for routine tasks like reproducibility checks, they catch fewer errors, propose weaker robustness checks, and ultimately move slower¹⁸. As scientists continue this shift towards AI, collective implications emerge. For instance, an analysis of 41 million research papers suggests that the adoption of AI tools and methods increases productivity, impact, and career outcomes for individuals — but at the expense of an overall narrowing of the scope of science¹⁹ (Figure 2).

Areas of consensus — and dissensus

Particularly after focusing a prior week on the benefits of AI, seminar discussion was marked by areas of consensus and dissensus^{7,8}.

In general, there was strong agreement that concerns about AI have evolved beyond hallucination proneness²⁰. Decreased vigilance and a collective narrowing of science received more attention

Top discussion points



Top resources



Your Brain on ChatGPT

54 essay writers under EEG: LLM users showed the weakest neural connectivity and least essay ownership.

Identified by Majed A. Arishi and Rui Jin



AI Tools Expand Impact, Contract Focus

41.3 million papers: AI-augmented scientists publish and are cited far more, while topics studied narrow.

Identified by Zhongchi Li and Sooyeon Park



AI-Assisted vs. AI-Led Research Teams

Human-only and AI-assisted teams reproduced over 90% of published results; AI-led teams managed 37%.

Identified by Paola J. Sullivan and Mariana G. Rebello

Figure 1: Overview of Week 3 of the Fall 2026 Collective [Artificial] Intelligence Seminar. Left: top discussion points, along with estimated level of agreement among seminar participants. Right: three most popular resources identified by seminar participants.

among seminar participants — with near-unanimous agreement that maximizing value from AI requires training, practice, and refinement²¹. Yet, there was clear disagreement about the magnitude of the issue. Does AI hinder more than help with critical thinking? In the midst of widespread concerns about the robustness of existing research practices²², to what extent should we aim for increases in productivity? Last, is a narrowing of the scope of science necessarily a bad tradeoff²³? After all, depression, addiction, cancer, dementia, climate change, and polarization are arguably enduring problems in need of sustained focus. Whether those problems are solved more quickly through increases in individual productivity, decreases in collective scope, and a faster overall science remains an open question.




Figure 2: Illustration of the narrowing of the scope of science, created by Zhongchi Li, Sooyeon Park, and Nicholas Coles.

Literature Review and Synthesis

Literature Review and Synthesis: Using AI

Covered in Week 4 of the seminar (September 10, 2026). Estimated date of arrival: September 20, 2026.

Literature Review and Synthesis: Critiquing AI

Covered in Week 5 of the seminar (September 17, 2026). Estimated date of arrival: September 27, 2026.

Data Collection and Generation

Data Collection and Generation: Using AI

Covered in Week 6 of the seminar (September 24, 2026). Estimated date of arrival: October 4, 2026.

Data Collection and Generation: Critiquing AI

Covered in Week 7 of the seminar (October 1, 2026). Estimated date of arrival: October 11, 2026.

Analysis and Modeling

Analysis and Modeling: Using AI

Covered in Week 8 of the seminar (October 8, 2026). Estimated date of arrival: October 18, 2026.

Analysis and Modeling: Critiquing AI

Covered in Week 9 of the seminar (October 15, 2026). Estimated date of arrival: October 25, 2026.

Writing and Communication

Writing and Communication: Using AI

Covered in Week 10 of the seminar (October 22, 2026). Estimated date of arrival: November 1, 2026.

Writing and Communication: Critiquing AI

Covered in Week 11 of the seminar (October 29, 2026). Estimated date of arrival: November 8, 2026.

Ethics and Outreach

Ethics and Outreach: Using AI

Covered in Week 12 of the seminar (November 5, 2026). Estimated date of arrival: November 15, 2026.

Ethics and Outreach: Critiquing AI

Covered in Week 13 of the seminar (November 12, 2026). Estimated date of arrival: November 22, 2026.

Personal AI Philosophy Statements

Covered in Week 14 of the seminar (November 19, 2026). Estimated date of arrival: November 29, 2026.

References

1. Coles, N. A., Hamlin, J. K., Sullivan, L. L., Parker, T. H. & Altschul, D. [Build up big-team science](#). *Nature* **601**, 505–507 (2022).
2. Coles, N. A., Goes Braga Takayanagi, J. F., Grant, G. L. & Basnight-Brown, D. M. [The rise, impact, and imbalances of big-team psychology](#). *Psychological Science* **37**, 287–297 (2026).
3. Radford, A., Narasimhan, K., Salimans, T. & Sutskever, I. *Improving Language Understanding by Generative Pre-Training*. https://cdn.openai.com/research-covers/language-unsupervised/language_understanding_paper.pdf (2018).
4. Brown, T. B., Mann, B., Ryder, N., Subbiah, M., *et al.* Language models are few-shot learners. *arXiv preprint* <https://doi.org/10.48550/arXiv.2005.14165> (2020) doi:10.48550/arXiv.2005.14165.
5. OpenAI. OpenAI GPT-5 system card. *arXiv preprint* <https://doi.org/10.48550/arXiv.2601.03267> (2025) doi:10.48550/arXiv.2601.03267.
6. OpenAI. The hugging face incident and the road ahead. <https://openai.com/index/hugging-face-incident-and-the-road-ahead/> (2026).
7. Coles, N. A. *et al.* [Team scientists should normalize disagreement](#). *Science* **384**, 1076–1077 (2024).
8. Dalkey, N. & Helmer, O. [An experimental application of the Delphi method to the use of experts](#). *Management Science* **9**, 458–467 (1963).
9. Baumgartner, H. A. *et al.* [How to build up big team science: A practical guide for large-scale collaborations](#). *Royal Society Open Science* **10**, 230235 (2023).
10. National Research Council. *Enhancing the Effectiveness of Team Science*. (National Academies Press, Washington, DC, 2015). doi:10.17226/19007.
11. Si, C., Yang, D. & Hashimoto, T. [Can LLMs generate novel research ideas? A large-scale human study with 100+ NLP researchers](#). in *International conference on learning representations* (eds Yue, Y., Garg, A., Peng, N., Sha, F. & Yu, R.) vol. 2025 94003–94092 (2025).
12. Tong, S., Mao, K., Huang, Z., Zhao, Y. & Peng, K. Automating psychological hypothesis generation with AI: When large language models meet causal graph. *Humanities and Social Sciences Communications* **11**, 896 (2024).
13. Schmitt, B., Hao, S., Pham, M. T. & Hofstetter, R. Generative grounded theory (GGT): Inductive theory building in the age of generative AI. *Journal of Consumer Psychology* <https://doi.org/10.1002/jcpy.70036> (2026) doi:10.1002/jcpy.70036.

14. Yamada, Y. *et al.* [Towards end-to-end automation of AI research](#). *Nature* **651**, 914–919 (2026).
15. Nature. Why AI cannot do good science without humans. *Nature* <https://www.nature.com/articles/d41586-026-01551-3> (2026).
16. Grigorev, A. [Claude Code wiped our production database with a Terraform command](#). (2026).
17. Kosmyrna, N. *et al.* Your brain on ChatGPT: Accumulation of cognitive debt when using an AI assistant for essay writing task. *arXiv preprint arXiv:2506.08872* <https://doi.org/10.48550/arXiv.2506.08872> (2025) doi:10.48550/arXiv.2506.08872.
18. Brodeur, A. *et al.* [AI-assisted teams outperform AI-led teams but not human-only teams in assessing research reproducibility in quantitative social science](#). *Proceedings of the National Academy of Sciences* **123**, e2524747123 (2026).
19. Hao, Q., Xu, F., Li, Y. & Evans, J. [Artificial intelligence tools expand scientists’ impact but contract science’s focus](#). *Nature* **649**, 1237–1243 (2026).
20. Topaz, M., Roguin, N., Gupta, P., Zhang, Z. & Peltonen, L.-M. [Fabricated citations: An audit across 2 · 5 million biomedical papers](#). *The Lancet* **407**, 1779–1781 (2026).
21. Southworth, J. *et al.* [Developing a model for AI across the curriculum: Transforming the higher education landscape via innovation in AI literacy](#). *Computers and Education: Artificial Intelligence* **4**, 100127 (2023).
22. Munafò, M. R. *et al.* [A manifesto for reproducible science](#). *Nature Human Behaviour* **1**, 0021 (2017).
23. Kuhn, T. S. The essential tension: Tradition and innovation in scientific research. in *The essential tension: Selected studies in scientific tradition and change* 225–239 (University of Chicago Press, Chicago, 1977). doi:10.7208/9780226217239-010.